

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [3]: 1 df_train = pd.read_csv('train.csv')
```

```
In [4]: 1 df_train
```

Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...	...	...	...	...	...	...	...	...	...	...	...	...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q

891 rows × 12 columns

## EDA

In [5]: 1 df\_train.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch          891 non-null    int64
8   Ticket          891 non-null    object
9   Fare           891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [6]: 1 df\_train.drop(['PassengerId', 'Name', 'Ticket'], axis = 1, inplace = True)

In [7]: 1 100\*df\_train.isna().sum()/df\_train.shape[0]## Percentage of Na values in our dataset

```
Out[7]: Survived      0.000000
Pclass      0.000000
Sex         0.000000
Age        19.865320
SibSp       0.000000
Parch       0.000000
Fare        0.000000
Cabin      77.104377
Embarked    0.224467
dtype: float64
```

In [8]: 1 df\_train.head()

Out[8]:


	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	0	3	male	22.0	1	0	7.2500	NaN	S
1	1	1	female	38.0	1	0	71.2833	C85	C
2	1	3	female	26.0	0	0	7.9250	NaN	S
3	1	1	female	35.0	1	0	53.1000	C123	S
4	0	3	male	35.0	0	0	8.0500	NaN	S

#### Questions to ask:

1. Analyse the distribution in X variable/Who were the passengers
2. Analyse the distribution in the Y variable/Understand the distribution in the Survived column
3. Analyse the relation between X and Y/ Find out which attributes are affecting the Y variable

```
In [9]: 1 ## Distribution in sex column  
2 sns.countplot(x = 'Sex', data = df_train)
```

Out[9]: <matplotlib.axes.\_subplots.AxesSubplot at 0x26e9558bdc8>




```
In [10]: 1 df_train['Sex'].value_counts()/df_train.shape[0]
```

Out[10]: male 0.647587  
female 0.352413  
Name: Sex, dtype: float64

```
In [11]: 1 ## Distribution in Pclass column  
2 sns.countplot(x = 'Pclass', data = df_train)
```

Out[11]: <matplotlib.axes.\_subplots.AxesSubplot at 0x26e956d67c8>




```
In [12]: 1 df_train['Pclass'].value_counts()/df_train.shape[0]
```

```
Out[12]: 3    0.551066  
1    0.242424  
2    0.206510  
Name: Pclass, dtype: float64
```

```
In [13]: 1 ## Distribution in Embarked column  
2 sns.countplot(x = 'Embarked', data = df_train)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x26e9573a788>
```




```
In [14]: 1 df_train['Embarked'].value_counts()/df_train.shape[0]
```

```
Out[14]: S    0.722783  
C    0.188552  
Q    0.086420  
Name: Embarked, dtype: float64
```

```
In [15]: 1 df_train['Age'].hist()
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95c01148>
```




```
In [16]: 1 df_train['Age'].describe()
```

```
Out[16]: count    714.000000  
mean       29.699118  
std        14.526497  
min         0.420000  
25%        20.125000  
50%        28.000000  
75%        38.000000  
max        80.000000  
Name: Age, dtype: float64
```

```
In [17]: 1 df_train['Fare'].hist()
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95a2de08>
```



```
In [18]: 1 df_train['Fare'].describe()
```


```
Out[18]: count    891.000000
mean         32.204208
std          49.693429
min           0.000000
25%           7.910400
50%          14.454200
75%          31.000000
max         512.329200
Name: Fare, dtype: float64
```

### ***Multivariate Analysis***




```
In [19]: 1 sns.countplot(x = 'Pclass',data = df_train,hue = 'Sex')  
2 ## homework: Find the numbers eg in pclass 3 what is the percentage of males vs females
```

Out[19]: <matplotlib.axes.\_subplots.AxesSubplot at 0x26e95b2fb48>




```
In [20]: 1 sns.countplot(x = 'Pclass',data = df_train,hue = 'Embarked')
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95c8f148>
```



```
In [21]: 1 sns.countplot(x = 'Embarked',data = df_train,hue = 'Pclass')
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95ce7d48>
```



```
In [22]: 1 df_train['Cabin'].isna().sum()/df_train.shape[0]
2 ## Depending on your domain and column importance if you have more than 25-30% na values in any particular column yo
```

```
Out[22]: 0.7710437710437711
```

```
In [23]: 1 for i in df_train['Embarked'].dropna().unique():
          2     print(f"Passangers Embarked from Port {i} : the distribution is \n {100*df_train[df_train['Embarked']==i]['Pclass']
```

Passangers Embarked from Port S : the distribution is

3 39.618406

2 18.406285

1 14.253648

Name: Pclass, dtype: float64

Passangers Embarked from Port C : the distribution is

1 9.539843

3 7.407407

2 1.907969

Name: Pclass, dtype: float64

Passangers Embarked from Port Q : the distribution is

3 8.080808

2 0.336700

1 0.224467

Name: Pclass, dtype: float64

```
In [24]: 1 df_train['Cabin']
```

Out[24]: 0 NaN

1 C85

2 NaN

3 C123

4 NaN

...

886 NaN

887 B42

888 NaN

889 C148

890 NaN

Name: Cabin, Length: 891, dtype: object

```
In [25]: 1 levels = []
          2
          3 levels = [i[0] for i in df_train['Cabin'].dropna()]
```

In [26]:

1 levels


```
'E',  
'E',  
'C',  
'C',  
'C',  
'F',  
'C',  
'E',  
'E',  
'B',  
'B',  
'D',  
'C',  
'B',  
'B',  
'D',  
'E',  
'B',  
'B',  
'D'
```

In [27]:

1 cabin = pd.DataFrame(levels,columns = ['Cabin'])

```
In [28]: 1 sns.countplot(x = 'Cabin', data = cabin)
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95d94548>
```



```
In [29]: 1 df_train['Alone'] = df_train['SibSp'] + df_train['Parch']
```

```
In [30]: 1 df_train['Alone'].loc[df_train['Alone']>0] = 'With Family'
        2 df_train['Alone'].loc[df_train['Alone']==0] = 'Alone'
```


C:\Users\yashm\anaconda3\lib\site-packages\pandas\core\indexing.py:1732: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self._setitem_single_block(indexer, value, name)
```

```
In [31]: 1 sns.countplot(x = 'Alone',data = df_train)
```


Out[31]: <matplotlib.axes.\_subplots.AxesSubplot at 0x26e95e23708>



**Analyse the distribution in the Y variable/Understand the distribution in the Survived column**

```
In [32]: 1 sns.countplot(x = 'Survived',data = df_train)
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95e671c8>
```




```
In [33]: 1 df_train['Survived'].value_counts()/df_train.shape[0]
```

```
Out[33]: 0    0.616162  
         1    0.383838  
         Name: Survived, dtype: float64
```



```
In [34]: 1 sns.countplot(x = 'Survived',data = df_train,hue = 'Sex')
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x26e95ebefc8>
```



In [35]: 1 df\_train

Out[35]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Alone
0	0	3	male	22.0	1	0	7.2500	NaN	S	With Family
1	1	1	female	38.0	1	0	71.2833	C85	C	With Family
2	1	3	female	26.0	0	0	7.9250	NaN	S	Alone
3	1	1	female	35.0	1	0	53.1000	C123	S	With Family
4	0	3	male	35.0	0	0	8.0500	NaN	S	Alone
...	...	...	...	...	...	...	...	...	...	...
886	0	2	male	27.0	0	0	13.0000	NaN	S	Alone
887	1	1	female	19.0	0	0	30.0000	B42	S	Alone
888	0	3	female	NaN	1	2	23.4500	NaN	S	With Family
889	1	1	male	26.0	0	0	30.0000	C148	C	Alone
890	0	3	male	32.0	0	0	7.7500	NaN	Q	Alone

891 rows × 10 columns

```
In [69]: 1 def Person( Passenger ):
2         Sex, Age = Passenger
3         if Age < 18:
4             return "child"
5         else:
6             return Sex
7
8
```


```
In [70]: 1 df_train['Person'] = df_train[['Sex', 'Age']].apply(Person, axis = 1)
```

```
In [71]: 1 df_train['Person']
```

```
Out[71]: 0      male
          1      female
          2      female
          3      female
          4      male
          ...
         886     male
         887     female
         888     female
         889     male
         890     male
          Name: Person, Length: 891, dtype: object
```


```
In [72]: 1 sns.countplot(x = 'Person',data = df_train)
```

```
Out[72]: <matplotlib.axes._subplots.AxesSubplot at 0x26e98cf2148>
```



```
In [73]: 1 sns.countplot(x = 'Person',data = df_train,hue = 'Survived')
```


```
Out[73]: <matplotlib.axes._subplots.AxesSubplot at 0x26e9a323808>
```



## Numerical variables vs categoricals


```
In [74]: 1 fig = sns.FacetGrid(df_train,hue = 'Person',aspect = 3)
2 fig.map(sns.kdeplot,'Age',shade = True)
3
4 oldest = df_train['Age'].max()
5
6 fig.set(xlim = (0,oldest))
7
8 fig.add_legend()
```

Out[74]: <seaborn.axisgrid.FacetGrid at 0x26e9a38aa48>



```
In [75]: 1 fig = sns.FacetGrid(df_train,hue = 'Pclass',aspect = 3)
2 fig.map(sns.kdeplot,'Age',shade = True)
3
4 oldest = df_train['Age'].max()
5
6 fig.set(xlim = (0,oldest))
7
8 fig.add_legend()
```


Out[75]: <seaborn.axisgrid.FacetGrid at 0x26e9a3df348>



**Conclusive graphs/Analyse the relation between X and Y/ Find out which attributes are affecting the Y variable**

```
In [77]: 1 # cat vs cat  
2  
3 sns.factorplot(x = 'Pclass',y = 'Survived',data = df_train)
```

Out[77]: <seaborn.axisgrid.FacetGrid at 0x26e9906a288>




```
In [78]: 1 sns.factorplot(x = 'Pclass',y = 'Survived',data = df_train,hue = 'Person')
```

C:\Users\yashm\anaconda3\lib\site-packages\seaborn\categorical.py:3669: UserWarning: The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed to `strip` in `catplot`.

```
warnings.warn(msg)
```

```
Out[78]: <seaborn.axisgrid.FacetGrid at 0x26e9a59dc08>
```






```
In [79]: 1 sns.factorplot(x = 'Pclass',y = 'Survived',data = df_train,hue = 'Alone')
```

C:\Users\yashm\anaconda3\lib\site-packages\seaborn\categorical.py:3669: UserWarning: The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed to `strip` in `catplot`.


```
warnings.warn(msg)
```

```
Out[79]: <seaborn.axisgrid.FacetGrid at 0x26e9a572f48>
```




```
In [80]: 1 # numerical vs cat/numerical vs numerical  
2  
3 sns.lmplot('Age', 'Survived', df_train)
```

Out[80]: <seaborn.axisgrid.FacetGrid at 0x26e9a5a9948>




```
In [83]: 1 sns.lmplot('Age', 'Survived', df_train, hue = 'Person', x_bins = [10, 20, 40, 60, 80])
```

```
Out[83]: <seaborn.axisgrid.FacetGrid at 0x26e9b795b08>
```



```
In [84]: 1 sns.lmplot('Fare', 'Survived', df_train)
```

```
Out[84]: <seaborn.axisgrid.FacetGrid at 0x26e9b81e608>
```



## Homework

Write down the inference that we have discussed and if you find anymore for each of these graphs

## Model Building

## Data Cleaning

```
In [85]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.svm import SVC
```

```
In [86]: 1 df_train.head()
```

Out[86]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked	Alone	Person
0	0	3	male	22.0	1	0	7.2500	NaN	S	With Family	male
1	1	1	female	38.0	1	0	71.2833	C85	C	With Family	female
2	1	3	female	26.0	0	0	7.9250	NaN	S	Alone	female
3	1	1	female	35.0	1	0	53.1000	C123	S	With Family	female
4	0	3	male	35.0	0	0	8.0500	NaN	S	Alone	male

```
In [88]: 1 df_train = df_train.drop(['Alone', 'Person', 'Cabin'],1)
```

C:\Users\yashm\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only  
 """Entry point for launching an IPython kernel.

```
In [89]: 1 df_train.isna().sum()
```

```
Out[89]: Survived      0
Pclass      0
Sex         0
Age        177
SibSp       0
Parch       0
Fare        0
Embarked    2
dtype: int64
```

```
In [90]: 1 df_train['Embarked'].value_counts()
```

```
Out[90]: S    644
         C    168
         Q     77
         Name: Embarked, dtype: int64
```

```
In [92]: 1 df_train['Embarked'].fillna('S',inplace = True)
```

```
In [95]: 1 df_train['Age'].fillna(df_train['Age'].mean(),inplace = True)
```

```
In [96]: 1 df_train
```

```
Out[96]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.000000	1	0	7.2500	S
1	1	1	female	38.000000	1	0	71.2833	C
2	1	3	female	26.000000	0	0	7.9250	S
3	1	1	female	35.000000	1	0	53.1000	S
4	0	3	male	35.000000	0	0	8.0500	S
...	...	...	...	...	...	...	...	...
886	0	2	male	27.000000	0	0	13.0000	S
887	1	1	female	19.000000	0	0	30.0000	S
888	0	3	female	29.699118	1	2	23.4500	S
889	1	1	male	26.000000	0	0	30.0000	C
890	0	3	male	32.000000	0	0	7.7500	Q

891 rows × 8 columns

```
In [98]: 1 df_train = pd.get_dummies(df_train,drop_first=True)
```

```
In [102]: 1 from sklearn.model_selection import train_test_split
```

```
In [106]: 1 X_train,X_test,y_train,y_test = train_test_split(df_train.drop('Survived',1),df_train['Survived']  
2                                     ,stratify=df_train['Survived'])
```

C:\Users\yashm\anaconda3\lib\site-packages\ipykernel\_launcher.py:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only  
 """Entry point for launching an IPython kernel.

```
In [100]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [107]: 1 scaler = StandardScaler()
```

```
In [109]: 1 X_train[['Age', 'Fare']] = scaler.fit_transform(X_train[['Age', 'Fare']])
```

```
In [111]: 1 X_test[['Age', 'Fare']] = scaler.transform(X_test[['Age', 'Fare']])
```

In [112]: 1 X\_train

Out[112]:

	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked_S
778	3	0.010684	0	0	-0.469543	1	1	0
272	2	0.852969	0	1	-0.239754	0	0	1
253	3	0.033110	1	0	-0.306175	1	0	1
767	3	0.070376	0	0	-0.469299	0	1	0
779	1	1.002034	0	1	3.507934	0	0	1
...	...	...	...	...	...	...	...	...
881	3	0.256708	0	0	-0.466450	1	0	1
125	3	-1.308478	1	0	-0.401086	1	0	0
21	2	0.331240	0	0	-0.366736	1	0	1
878	3	0.010684	0	0	-0.466450	1	0	1
707	1	0.927501	0	0	-0.107155	1	0	1

668 rows × 8 columns

In [113]: 1 y\_train

Out[113]:

778	0
272	1
253	0
767	0
779	1
...	..
881	0
125	1
21	1
878	0
707	1

Name: Survived, Length: 668, dtype: int64



In [119]:

```

1  ##
2  from sklearn.metrics import classification_report
3  model = LogisticRegression()
4  model.fit(X_train,y_train)
5  print("Results of The Logistic Regression Model is : \n")
6  print(classification_report(model.predict(X_test),y_test))
7  print("-----*****")
8
9  model = KNeighborsClassifier()
10 model.fit(X_train,y_train)
11 print("Results of The Knn is : \n")
12 print(classification_report(model.predict(X_test),y_test))
13 print("-----*****")
14
15 model = DecisionTreeClassifier()
16 model.fit(X_train,y_train)
17 print("Results of The Decision Tree is : \n")
18 print(classification_report(model.predict(X_test),y_test))
19 print("-----*****")
20
21 model = RandomForestClassifier()
22 model.fit(X_train,y_train)
23 print("Results of The Random Forest is : \n")
24 print(classification_report(model.predict(X_test),y_test))
25 print("-----*****")
26
27 model = SVC()
28 model.fit(X_train,y_train)
29 print("Results of The Support Vector Machine is : \n")
30 print(classification_report(model.predict(X_test),y_test))
31 print("-----*****")

```

Results of The Logistic Regression Model is :

	precision	recall	f1-score	support
0	0.83	0.79	0.81	144
1	0.65	0.71	0.68	79
accuracy			0.76	223

macro avg	0.74	0.75	0.75	223
weighted avg	0.77	0.76	0.76	223

-----\*\*\*\*\*-----

Results of The Knn is :

	precision	recall	f1-score	support
0	0.82	0.80	0.81	142
1	0.66	0.70	0.68	81
accuracy			0.76	223
macro avg	0.74	0.75	0.75	223
weighted avg	0.77	0.76	0.76	223

-----\*\*\*\*\*-----

Results of The Decision Tree is :

	precision	recall	f1-score	support
0	0.77	0.82	0.79	130
1	0.72	0.67	0.69	93
accuracy			0.75	223
macro avg	0.75	0.74	0.74	223
weighted avg	0.75	0.75	0.75	223

-----\*\*\*\*\*-----

Results of The Random Forest is :

	precision	recall	f1-score	support
0	0.82	0.82	0.82	136
1	0.72	0.71	0.72	87
accuracy			0.78	223
macro avg	0.77	0.77	0.77	223
weighted avg	0.78	0.78	0.78	223

-----\*\*\*\*\*-----

Results of The Support Vector Machine is :

	precision	recall	f1-score	support
0	0.89	0.80	0.84	153
1	0.64	0.79	0.71	70
accuracy			0.79	223
macro avg	0.77	0.79	0.77	223
weighted avg	0.81	0.79	0.80	223

-----\*\*\*\*\*-----

In [ ]:

1